

# Loop Invariant Code Motion: Optimization and Evaluation

Chen, Hao-Wen

Department of Computer Science and Information Engineering  
National Taiwan University  
Email: b11902156@ntu.edu.tw

**Abstract**—Context-Sensitive Function Inlining (CSFI) is an advanced compiler optimization that selectively inlines function calls based on their usage context. This report provides a detailed overview of the algorithm design, implementation methodology, and correctness proof for CSFI within the LLVM framework.

## I. ALGORITHM DESIGN

### A. Optimization Strategy

Context-Sensitive Function Inlining improves runtime efficiency by reducing function call overhead and enabling further optimizations. Unlike naive inlining, CSFI considers the context of each call site, such as argument values, caller characteristics, and call frequency.

### B. Theoretical Analysis

Theoretical benefits include reduced function call overhead, increased instruction locality, and enhanced opportunities for constant folding and dead code elimination. However, these benefits must be weighed against the potential for increased code size (code bloat) and instruction cache pressure.

### C. Implementation Considerations

Key considerations include:

- Profiling to identify hot call sites.
- Analysis of argument-dependent behaviors.
- Balancing inlining benefits against code size growth.
- Handling recursive functions and indirect calls.

### D. Code Example

Consider the following code snippet:

```
int add(int x, int y) {
    return x + y;
}

void compute() {
    for (int i = 0; i < 100; i++) {
        int result = add(i, 10);
        process(result);
    }
}
```

After applying CSFI:

```
void compute() {
    for (int i = 0; i < 100; i++) {
        // Inlined version of add
        int result = i + 10;
        process(result);
    }
}
```

## II. IMPLEMENTATION DETAILS

### A. LLVM Pass Methodology

The CSFI pass in LLVM analyzes the call graph, profiling data, and call site characteristics to determine which functions to inline. It then performs inlining selectively, ensuring that the optimization aligns with the overall performance goals.

### B. Key Data Structures

- CallGraph: Represents the function call relationships.
- FunctionInfo: Stores metadata about functions, such as their size, frequency, and side effects.
- ProfileInfo: Contains runtime profiling data.

### C. Integration with LLVM Pipeline

CSFI integrates with LLVM's mid-level optimizations. It is typically executed after profiling-guided optimizations (PGO) and before interprocedural optimizations.

### D. Edge Case Handling

Special cases include:

- Recursive functions, where inlining may lead to infinite expansion.
- Indirect calls, requiring dynamic resolution or conservative inlining.
- Functions with significant side effects or large code size.

## III. EXPERIMENTAL EVALUATION

### A. Test Cases and Benchmarks

CSFI is evaluated using both synthetic and real-world benchmarks. Synthetic test cases focus on loop-heavy workloads, argument sensitivity, and code bloat scenarios. Real-world benchmarks include programs from the SPEC CPU suite and the LLVM test suite to measure performance improvements and code size changes. Profiling-driven tests further assess the impact on hot and cold paths. Results are compared against naive inlining and manual optimization to validate efficiency and scalability.

## B. Simple Inlining

```
int multiply(int a, int b) {
    return a * b;
}

void test() {
    int result = multiply(2, 3);
    printf("%d\n", result);
}
```

After inlining:

```
void test() {
    int result = 2 * 3;
    printf("%d\n", result);
}
```

## C. Recursive Function Handling

```
int factorial(int n) {
    if (n <= 1) return 1;
    return n * factorial(n - 1);
}

void compute_factorial() {
    int result = factorial(5);
    printf("%d\n", result);
}
```

Inlining is skipped for recursive calls to avoid infinite expansion.

## D. Context-Sensitive Decisions

```
int compute_offset(int base, int offset) {
    return base + offset;
}

void use_offset() {
    int x = compute_offset(10, 5);
    int y = compute_offset(20, 7);
    printf("%d %d\n", x, y);
}
```

After CSFI:

```
void use_offset() {
    int x = 10 + 5;
    int y = 20 + 7;
    printf("%d %d\n", x, y);
}
```

## E. Mixed Function Calls

In this test case, two different functions (`add` and `subtract`) are called within the same context:

```
int add(int x, int y) {
    return x + y;
}

int subtract(int x, int y) {
    return x - y;
}

void test_mixed_calls() {
    int sum = add(15, 5);
    int diff = subtract(20, 10);
    printf("%d %d\n", sum, diff);
}
```

CSFI selectively inlines these calls based on the context, potentially reducing overhead for frequently called functions.

*1) Test Case 5: Edge Case - Large Loop:* This test case demonstrates the challenge of inlining when the function is used in a large loop, where the performance gain of inlining is balanced against code size growth:

```
int increment(int x) {
    return x + 1;
}

void test_large_loop() {
    int sum = 0;
    for (int i = 0; i < 10000; i++) {
        sum += increment(i);
    }
    printf("%d\n", sum);
}
```

Inlining the `increment` function might increase code size, so CSFI must carefully consider whether the optimization is beneficial.

## IV. CORRECTNESS PROOF

### A. Invariant Preservation

CSFI preserves correctness by ensuring that:

- The inlined function's semantics match the original call.
- Side effects, such as memory writes, are accurately reproduced.
- All control flow paths are equivalent.

### B. Edge Case Analysis

Edge cases such as exception handling, dynamic memory allocation, and floating-point precision are handled by careful analysis of inlined code.

### C. Testing Methodology

Tests include:

- Synthetic benchmarks for common patterns.
- Real-world applications to evaluate scalability.
- Validation using LLVM's `opt` and `clang` tools.

## V. CONCLUSION

Context-Sensitive Function Inlining enhances performance by tailoring inlining decisions to specific call sites. Its integration into LLVM ensures robust and scalable optimization.